


Stochastic Gradient Descent (SGD)



Dataset D



Trained Model

Repeat ..

Select a random subset
(batch) from D



Compute gradients $\nabla L(\theta)$
on the batch



Update
 $\theta' = \theta - \eta * \nabla L(\theta)$

.. until model convergence (i.e., many epochs)

Differential Privacy (Weaker notion)

X : The data universe.

$D \subset X$: The dataset (one element per person)